

Your grade: 90%

Your latest: 90% • Your highest: 90% • To pass you need at least 80%. We keep your highest score.

Next item →

1. Suppose you learn a word embedding for a vocabulary of 10000 words. Then the embedding vectors could be 10000 dimensional, so as to capture the full range of variation and meaning in those words.

1 / 1 point

- ☐ True  
☒ False

✓ **Correct**  
The dimension of word vectors is usually smaller than the size of the vocabulary. Most common sizes for word vectors range between 50 and 1000.

2. True/False: t-SNE is a linear transformation that allows us to solve analogies on word vectors.

1 / 1 point

- ☒ False  
☐ True

✓ **Correct**  
tr-SNE is a non-linear dimensionality reduction technique.

3. Suppose you download a pre-trained word embedding which has been trained on a huge corpus of text. You then use this word embedding to train an RNN for a language task of recognizing if someone is happy from a short snippet of text, using a small training set.

1 / 1 point

| x (input text)        | y (happy?) |
|-----------------------|------------|
| Having a great time!  | 1          |
| I'm sad it's raining. | 0          |
| I'm feeling awesome!  | 1          |

Even if the word "wonderful" does not appear in your small training set, what label might be reasonably expected for the input text "I feel wonderful!"?

- ☐ y=0  
☒ y=1

✓ **Correct**  
Yes, word vectors empower your model with an incredible ability to generalize. The vector for "wonderful" would contain a negative/unhappy connotation which will probably make your model classify the sentence as a "1".

4. Which of these equations do you think should hold for a good word embedding? (Check all that apply)

1 / 1 point

☒  $e_{man} - e_{woman} \approx e_{king} - e_{queen}$

✓ **Correct**  
The order of words is correct in this analogy.

☐  $e_{man} - e_{king} \approx e_{queen} - e_{woman}$

☒  $e_{man} - e_{king} \approx e_{woman} - e_{queen}$

✓ **Correct**  
The order of words is correct in this analogy.

☐  $e_{man} - e_{woman} \approx e_{queen} - e_{king}$

5. Let  $A$  be an embedding matrix, and let  $o_{4567}$  be a one-hot vector corresponding to word 4567. Then to get the embedding of word 4567, why don't we call  $A * o_{4567}$  in Python?

1 / 1 point

- ☐ The correct formula is  $A^T * o_{4567}$ .  
☐ None of the answers are correct: calling the Python snippet as described above is fine.  
☐ This doesn't handle unknown words (<UNK>).  
☒ It is computationally wasteful.

✓ **Correct**  
Yes, the element-wise multiplication will be extremely inefficient.

6. When learning word embeddings, words are automatically generated along with the surrounding words.

1 / 1 point

- ☐ True  
☒ False



Correct

We pick a given word and try to predict its surrounding words or vice versa.

7. True/False: In the word2vec algorithm, you estimate  $P(t | c)$ , where  $t$  is the target word and  $c$  is a context word.  $t$  and  $c$  are chosen from the training set to be nearby words.

1 / 1 point

- ☒ True  
☐ False



Correct

Yes,  $t$  and  $c$  are chosen from the training set to be nearby words.

8. Suppose you have a 10000 word vocabulary, and are learning 100-dimensional word embeddings. The word2vec model uses the following softmax function:

1 / 1 point

$$P(t | c) = \frac{e^{\theta_t^T e_c}}{\sum_{t=1}^{10000} e^{\theta_t^T e_c}}$$

True/False: After training, we should expect  $\theta_t$  to be very close to  $e_c$  when  $t$  and  $c$  are the same word.

- ☐ True  
☒ False



Correct

To review this concept watch the *Word2Vec* lecture.

9. Suppose you have a 10000 word vocabulary, and are learning 500-dimensional word embeddings. The GloVe model minimizes this objective:

0 / 1 point

$$\min \sum_{i=1}^{10,000} \sum_{j=1}^{10,000} f(X_{ij})(\theta_i^T e_j + b_i + b_j' - \log X_{ij})^2$$

True/False:  $\theta_i$  and  $e_j$  should be initialized to 0 at the beginning of training.

- ☒ True  
☐ False



Incorrect

No,  $\theta_i$  and  $e_j$  should be initialized randomly at the beginning of training.

10. You have trained word embeddings using a text dataset of  $t_1$  words. You are considering using these word embeddings for a language task, for which you have a separate labeled dataset of  $t_2$  words. Keeping in mind that using word embeddings is a form of transfer learning, under which of these circumstances would you expect the word embeddings to be helpful?

1 / 1 point

- ☒ When  $t_1$  is larger than  $t_2$   
☐ When  $t_1$  is equal to  $t_2$   
☐ When  $t_1$  is smaller than  $t_2$



Correct

Transfer embeddings to new tasks with smaller training sets.